

TempusCaseStudy_RohanDhamdhare

March 13, 2017

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

#Loading Data (530 x 16563)
df = pd.read_table('C:\Users\dcc\Desktop\Tempus\takehome1.txt')
Data = df.values

#Split the Feature data and Ground Truth data
Xdata = Data[:,1:].copy()    # Copying to avoid damage to actual data
GTdata = Data[:,0].copy()

In [ ]: '''
Given the high number of Dimensions(16563), usage of some dimensionality
reduction technique is a must, to remove highly correlated features.
'''

In [2]: #Checking Proportion of Variance for various value of
# components in PCA transformed data

from sklearn.preprocessing import scale
from sklearn.decomposition import PCA

Data = scale(Data)          # Normalize data

#Number of Output Components for PCA
components = [100, 300, 400, 500, 700]
i=1
for component in components:
    pca = PCA(n_components=component, whiten= True)
    pca.fit(Data)

    #The amount of variance that each PC explains
    var= pca.explained_variance_ratio_

    #Cumulative Variance explains
    var1=np.cumsum(np.round(pca.explained_variance_ratio_, decimals=4)*100)
```


```



plt.figure(i)
plt.plot(var1)
plt.xlabel('Principal Components = %d' % component)
plt.ylabel('Expained Varaince (%)')
plt.title('Comparison')
plt.show()
i = i+1



```

C:\Users\dcc\Anaconda2\lib\site-packages\sklearn\utils\validation.py:429: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using y.ravel().

warnings.warn(msg, _DataConversionWarning)







```
In [ ]: '''
        -The above scree plots display the PCA Components needed to explain
        some percent variance in Data .
        -As it is evident from the plots, almost 100% variance in the data
        is explained by about 500 PCA components.
        -Thus the given 16563 variables are highly correlated and we can easily
        reduce the number of dimensions using PCA.
        - Almost 98% variance in data is explained by 400 PCA components.
        -To avoid overfitting to the data, i use only 380 PCA components to
        predict the 'response' variable.

        '''
```

```
In [ ]: '''
        -Given variable for prediction, 'response' is binary in nature.
        Thus, this is a 2-class classification problem.
        -After careful observation of data and Ground Truth(GT) values,
        it can be easily detected that there is a class imbalance, with 76%
        of the GT values being 0.
        -To avoid this problem, we divide our data into Training and Test sets
        using StratifiedKfold cross validation.
        -This method splits the data, such that the ratio of number of
        elements of each class is equal in set of Data.
        -I split the data as follows: Train set - 75%, Test set - 25%
        -Further, i divide the Train set data as follows -
        Train_actual = 75% Train set, Validation = 25% Train set.

        '''
```

```
In [ ]: '''
        -As explained in the case study document, no model is a perfect model.
        -So i provide a comparative analysis of 4 well known classification
        models:
        1] Linear SVM, 2] Logistic Regression, 3] Decision tree,
        4] Neural Nets.

        -As observed before, there is a class imbalance in data, with 76% of
        the GT values being 0.
        -As there is class imbalance in GT values, Accuracy is not the best
        metric to judge a classifier.
        -With predicting all zeros, a classifier can achieve an accuracy
        of 76%

        -Thus i use F-Measure of the classifier as a metric to judge the
        best classifier for the given data.
            
$$F\text{-Measure} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

        -I have displayed the following metrics :
```

*-Accuracy
-Precision
-Recall
-F-Measure*

Finally, i choose the best Classifier based of the F-measure.

'''

```
In [3]: from sklearn.pipeline import Pipeline
        from sklearn.model_selection import StratifiedKFold, KFold

        #Importing Classifiers
        from sklearn import linear_model
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.svm import SVC
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.neural_network import MLPClassifier

        #Importing Classification Metrics
        from sklearn.metrics import precision_score
        from sklearn.metrics import recall_score
        from sklearn.metrics import f1_score
        from sklearn.metrics import accuracy_score

        names = ["Logistic Regression", "Linear SVM", "Decision Tree", "NNet"]
        classifiers = [
            linear_model.LogisticRegression(C= 0.1),
            SVC(kernel="linear", C=0.025),
            DecisionTreeClassifier(max_depth=5),
            MLPClassifier(alpha=0.9,hidden_layer_sizes=(20,20))
        ]

        best_classifier = {}

        for name, classifier in zip(names, classifiers):

            #initialize metrics
            best_precision = 0.00
            best_recall = 0.00
            best_Fmeasure = 0.00
            best_accuracy = 0.00

            # Divide Data into Train and Test sets
            skf = StratifiedKFold(n_splits = 4)
```

```

    for Mtrain, test in skf.split(Xdata, GTdata):
        Xtrain = Xdata[Mtrain, :].copy()
        Xtest = Xdata[test, :].copy()
        GTtrain = GTdata[Mtrain]
        GTtest = GTdata[test]

# Divide Train data into Train_actual and Validation
kf = KFold(n_splits = 4)
for train, val in kf.split(Xtrain, GTtrain):
    Xtrain_actual = Xtrain[train, :].copy()
    Xval = Xtrain[val, :].copy()
    GTtrain_actual = GTtrain[train]
    GTval = GTtrain[val]

# PCA

pca = PCA(n_components=380, whiten=True)

# Pipeline Model

pipe = Pipeline([('pca', pca),
                  ('classifier', classifier)])
pipe.fit(Xtrain, GTtrain)
predicted = pipe.predict(Xtest)

# Metrics

precision = precision_score(GTtest, predicted)
if precision > best_precision: best_precision=precision
recall = recall_score(GTtest, predicted)
if recall > best_recall: best_recall = recall
accuracy = accuracy_score(GTtest, predicted)
if accuracy > best_accuracy: best_accuracy = accuracy
Fmeasure = f1_score(GTtest, predicted)
if Fmeasure > best_Fmeasure: best_Fmeasure = Fmeasure
best_classifier[best_Fmeasure] = name

print ("Classifier: " + str(name))
print ("Accuracy: " + str(round(best_accuracy, 3)))
print ("Precision: " + str(round(best_precision, 3)) + "\nRecall: " + str(round(best_recall, 3)))
print ("F-Measure = " + str(round(best_Fmeasure, 3)) + "\n")

best = best_classifier.keys()
Label = best_classifier.values()
r = np.arange(1, len(best)+1, 1)
best_Fmeasure = max(best)
Fmeasure_of_best = best_classifier.keys()[best_classifier.values().index(best_Fmeasure)]

print ("Best Model is: " + str(best_classifier[best_Fmeasure]))
print ("Fmeasure: " + str(round(Fmeasure_of_best, 3)))

```

```
Classifier: Logistic Regression
Accuracy: 0.887
Precision:0.864
Recall: 0.613
F-Measure = 0.717
```

```
Classifier: Linear SVM
Accuracy: 0.88
Precision:0.857
Recall: 0.581
F-Measure = 0.692
```

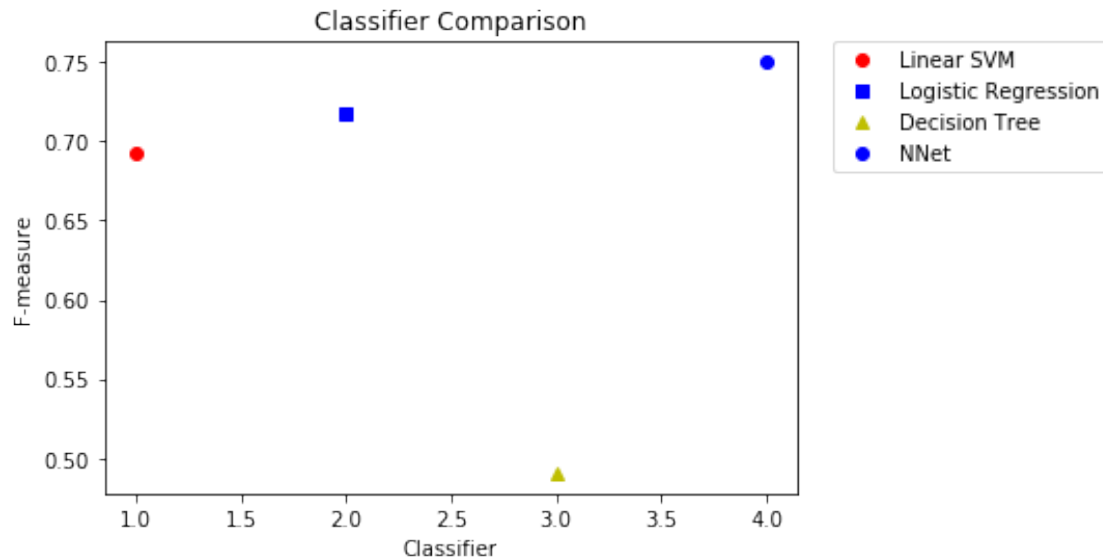
```
Classifier: Decision Tree
Accuracy: 0.805
Precision:0.857
Recall: 0.419
F-Measure = 0.491
```

```
C:\Users\dcc\Anaconda2\lib\site-packages\sklearn\neural_network\multilayer_perceptron.py:379: ConvergenceWarning:
% (), ConvergenceWarning)
```

```
Classifier: NNet
Accuracy: 0.895
Precision:1.0
Recall: 0.677
F-Measure = 0.75
```

```
Best Model is: NNet
Fmeasure: 0.75
```

```
In [4]: #Plotting the Classifier Comparison
import itertools
marker = itertools.cycle(('ro', 'bs', 'y^', 'bo', 'y*', 'b^'))
fig, ax = plt.subplots()
for i in range(len(r)):
    ax.plot(r[i], best[i], marker.next(), label=Label[i])
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
plt.xlabel('Classifier')
plt.ylabel('F-measure')
plt.title('Classifier Comparison')
plt.show()
```

```
In [ ]: '''
- After careful observation, I had noted that 76% Ground Truth
data is from class 0.
- So even predicting all zeros will give me an accuracy of
76% in prediction.
- If any model that gives an accuracy less than 76%, is actually
harming the prediction model.
- Using my machine learning knowledge, i used a few classification
models for prediction. All of them give an Accuracy > 77%.
```

After observing the above metric based results and F-measure based plot for the different models, we can infer the following:

```
F-Measure based Classifier ranking(1=worst):
4)Neural Net (0.75)
3)Logistic Regression(0.717)
2)Linear SVM(0.692)
1)Decision Trees(0.491)
```

```
Accuracy based Classifier ranking(1=worst):
4)Neural Net (89.5%)
3)Logistic Regression(88.7%)
2)Linear SVM (88%)
1)Decision Tree(80.5%)
```

*** All these classifier metrics are naive, and their no*

hyperparameter tuning has been done.

-Except for Decision trees, all other classifiers have a comparable F-measure.

-Thus, for me, the best model would be the one which would give the best accuracy for a comparable F-measure.

From the above rankings, a model with maximum value after summing the individual ranks would be the best model.

Here, undisputedly, it is Neural Network.

But if a simpler Model is required, Logistic Regression is a Good bet!

Thus here the best model, Neural Net, has a predicts the data with almost 90% accuracy.

'''